

# Generative Adversarial Networks (GANs)

**Quan Zhang, Ph.D.**

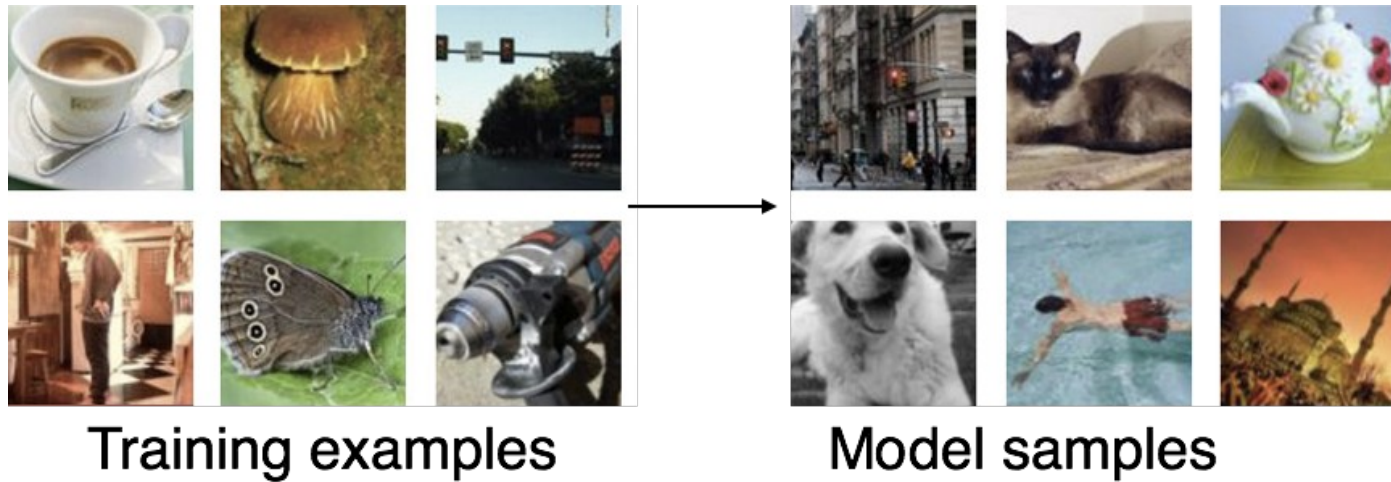

- Goodfellow, Ian; Pouget-Abadie, Jean; Mirza, Mehdi; Xu, Bing; Warde-Farley, David; Ozair, Sherjil; Courville, Aaron; Bengio, Yoshua (2014). *Generative Adversarial Networks*. Proceedings of the International Conference on Neural Information Processing Systems (NIPS 2014). pp. 2672–2680.

# Probabilistic Generative Models




**Density Estimation**

# Synthesizing Examples From Probabilistic Generative Model



# Maximum Likelihood Estimation



$$\theta^* = \arg \max_{\theta} \mathbb{E}_{x \sim p_{\text{data}}} \log p_{\text{model}}(x \mid \theta)$$

# Density function

## Explicit and analytical

- e.g., Gaussian
- can sample directly from model

## Explicit and approximate

- e.g., Gaussian mixture models, zero-inflated negative binomial models
- can estimate probability by running Markov chain monte carlo

## Implicit

- GAN
- can't estimate probability but can draw from distribution with given probability

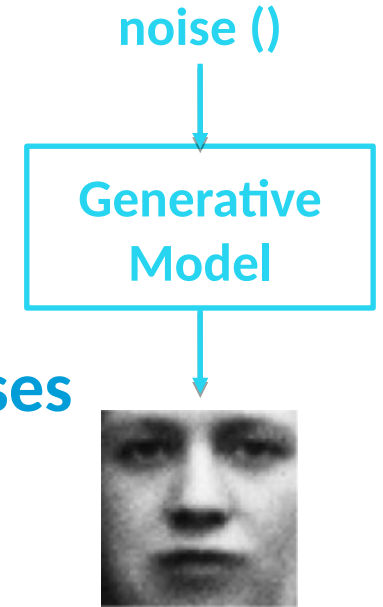
# Generative Model

How to make it generate different samples each time it is run?

- input to model is noise

Generative model as a neural network

- Generate random samples by computing a transformation of noises
- differentiable
- does not have to be invertible
- typically has very high dimensionality



# Discriminative Model

Think of it as a critic

- a good critic can tell real from fake

Discriminative model as a neural net

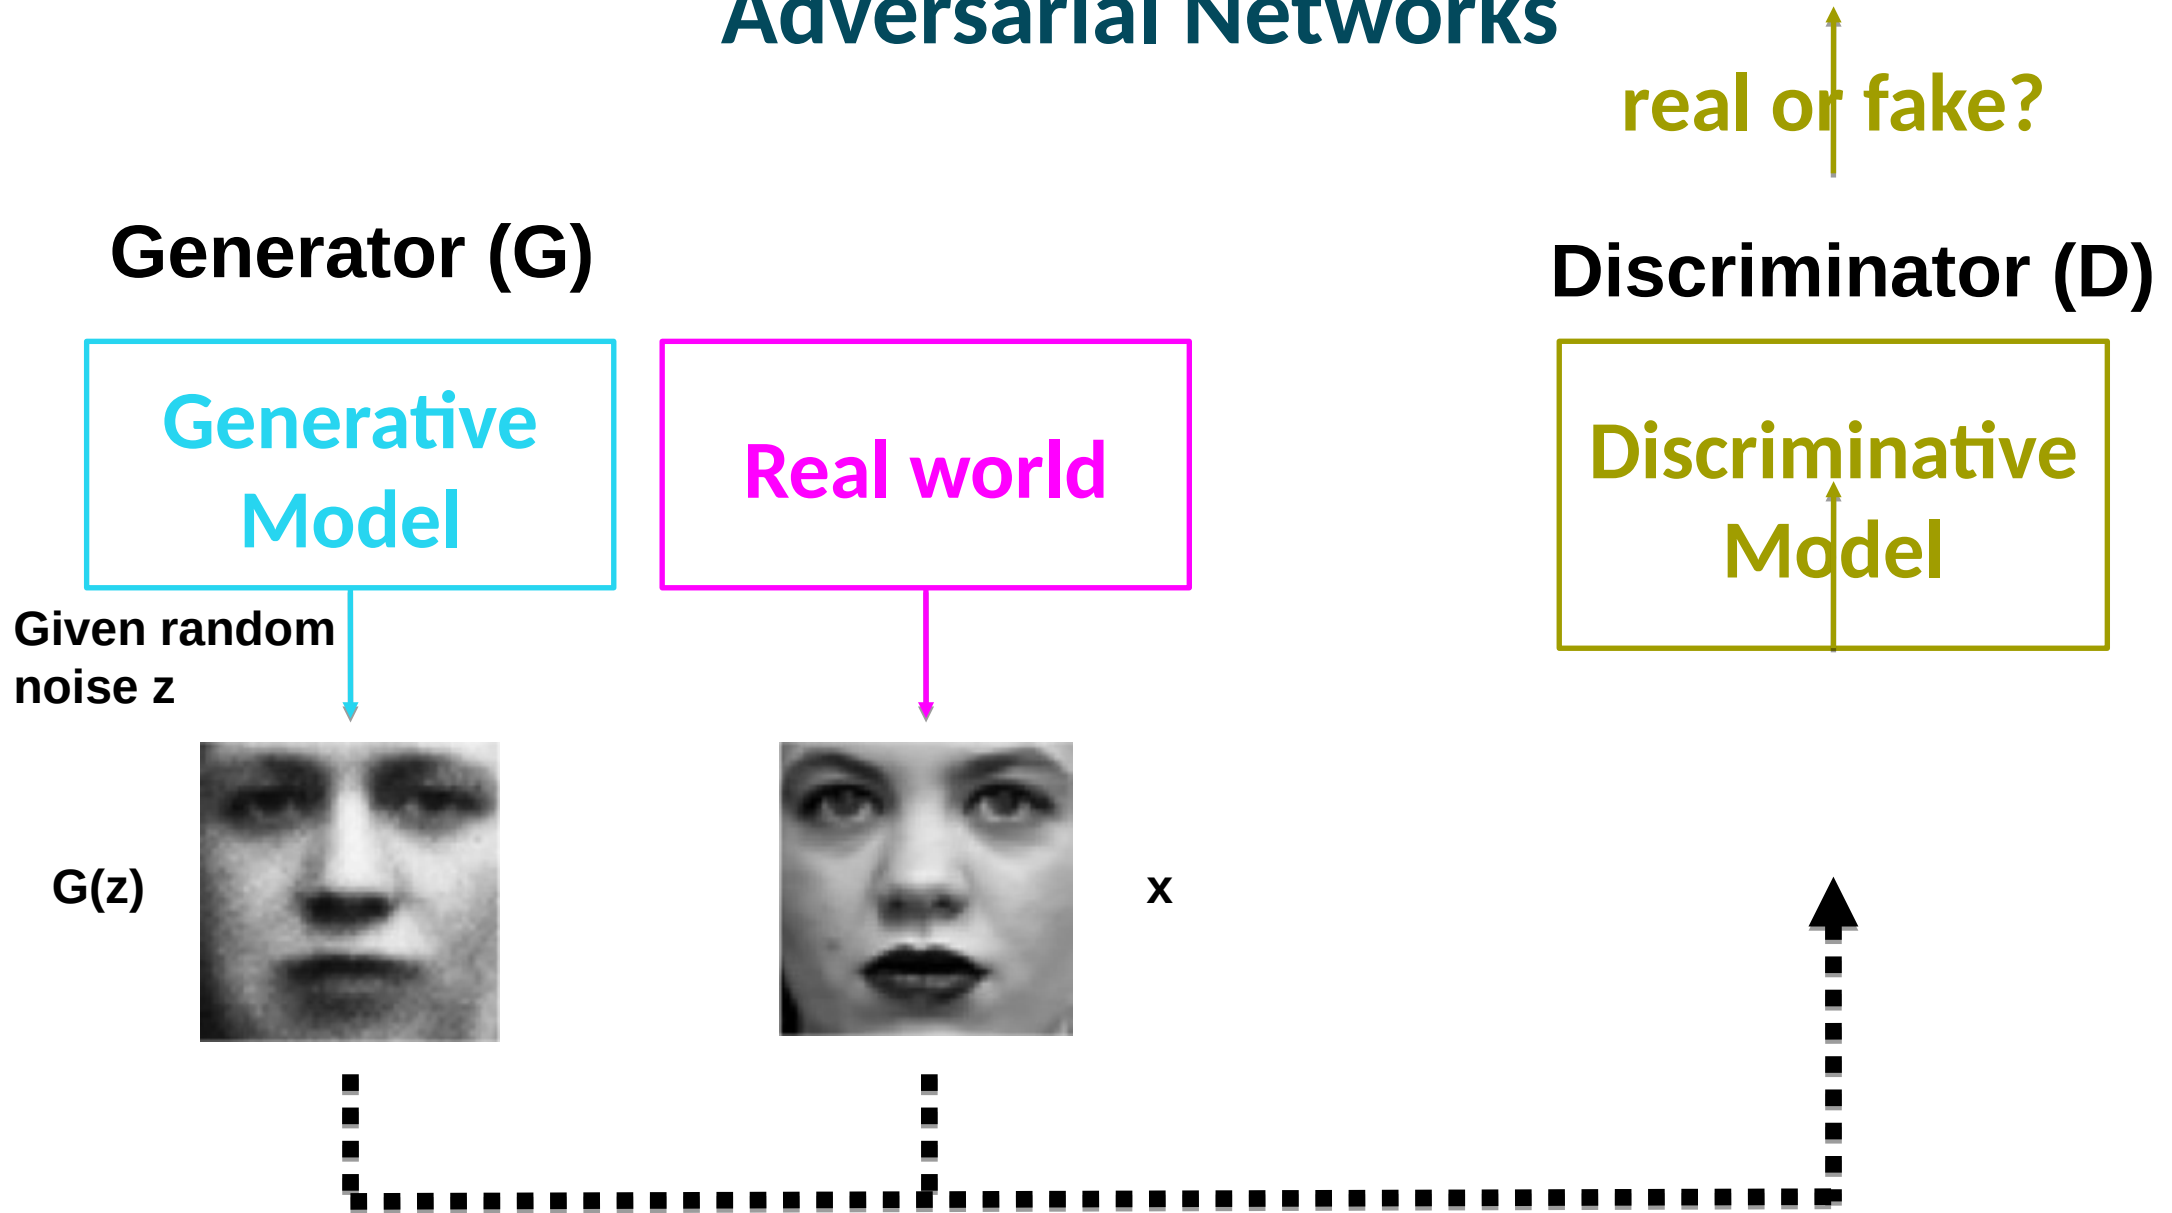
- differentiable
- computes  $\sigma$ , with value 1 if real, 0 if fake

real or fake?





# Adversarial Networks



# Training Procedure: Basic Idea

G tries to fool D


D tries not to be fooled

Models are trained simultaneously

- As G gets better, D has a more challenging task
- As D gets better, G has a more challenging task

Ultimately, we don't care about the D

- Its role is to force G to work harder



# Loss Functions

Loss function for D:  $x$  is real data,  $G(z)$  is generated data

- maximize the likelihood that model says 'real' to samples from the world and 'fake' to generated samples

$$\max \log D(x) + \log(1 - D(G(z)))$$

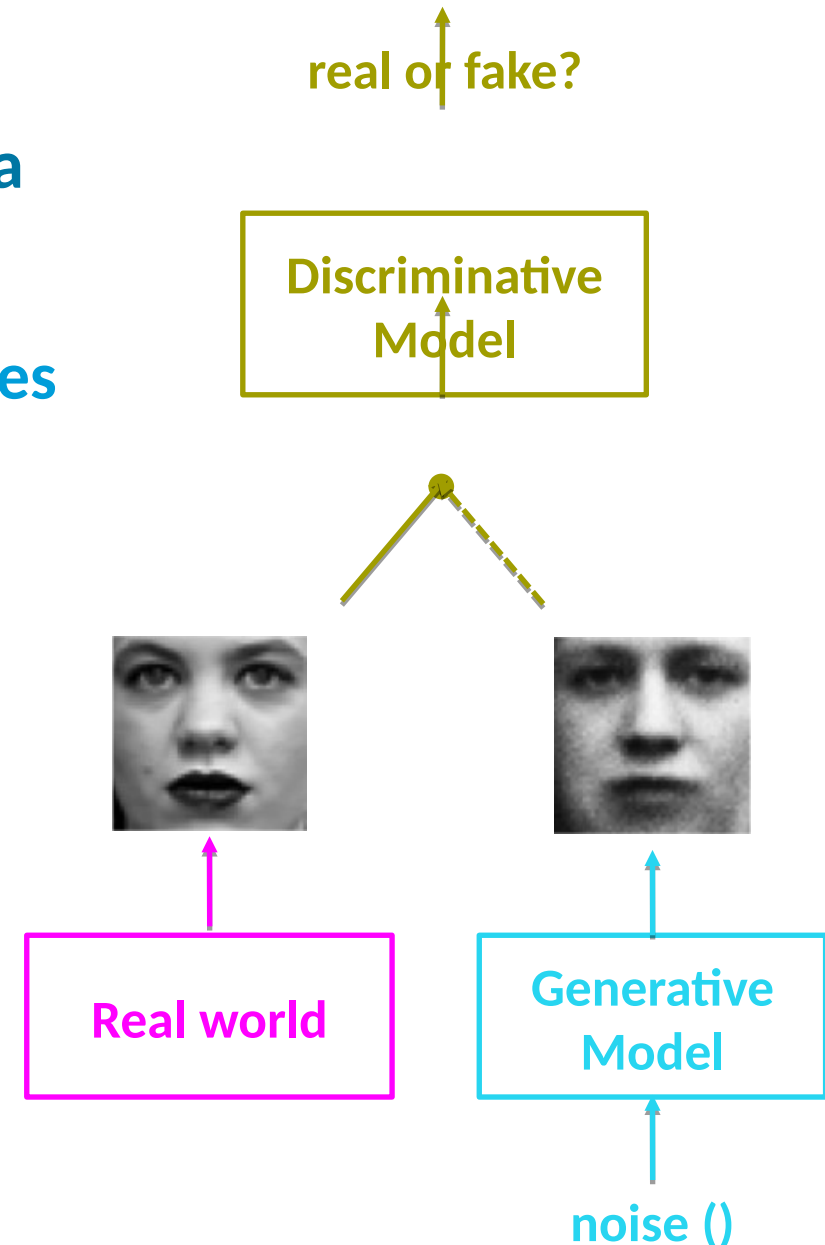
What should the loss function be for G?

$$\min \log D(x) + \log(1 - D(G(z)))$$

But because first term doesn't matter for G (why?),

$$\min \log(1 - D(G(z)))$$

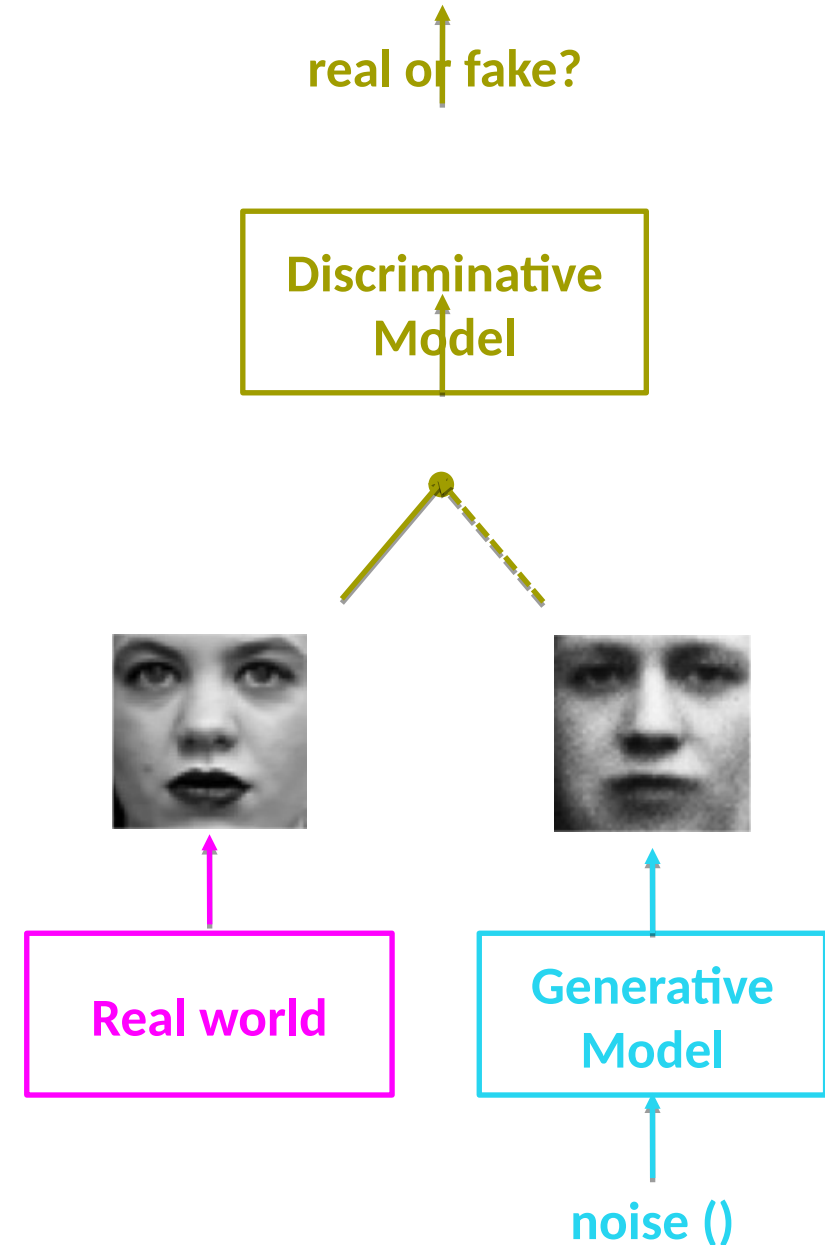
Known as a minimax procedure



# Training Procedure

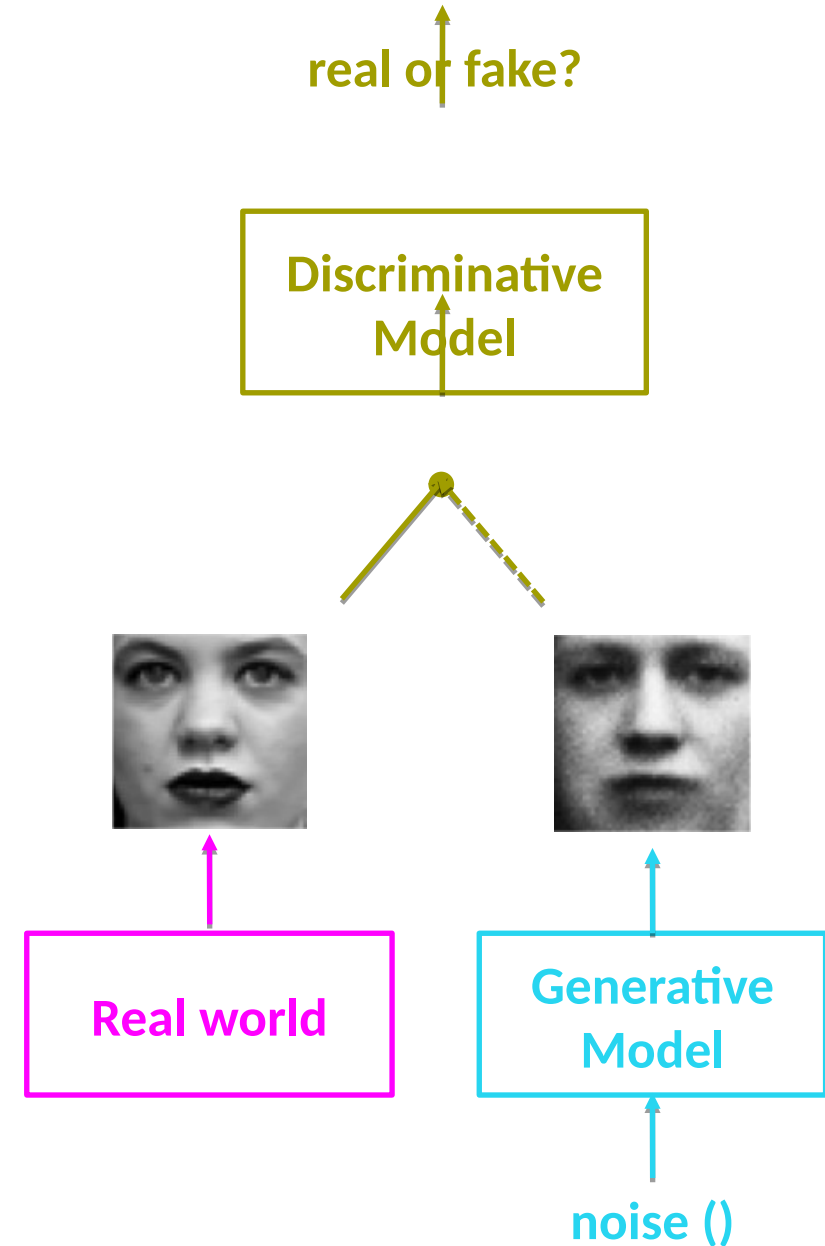

Train both models simultaneously via stochastic gradient descent using minibatches consisting of

- some generated samples
- some real-world samples





# The Discriminator Has a Straightforward Task



D has learned when



# Cherry Picked Results





# Problems With Counting





# Problems With Global Structure





# Visually-Aware Fashion Recommendation and Design

## With GANs (Kang et al., 2017)

Recommender systems predict how much a particular user will like a particular item


- Can predict based on features of item (e.g., movie director, dress length)
- Can also predict directly from images

Twist here is that instead of predicting from a predefined set, *generate images* that would be liked.

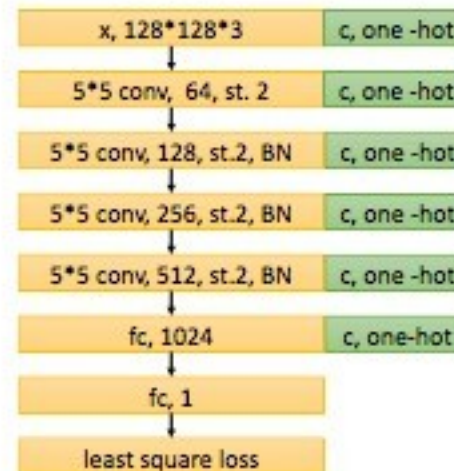


# Visually-Aware Fashion Recommendation and Design With GANs (Kang et al., 2017)

Use class-conditional generator and discriminator



(a) Generator  $G(z, c)$



(b) Discriminator  $D(x, c)$

# Visually-Aware Fashion Recommendation and Design With GANs (Kang et al., 2017)

## Optimize with GAN

- find latent representation that obtains the highest recommendation score
- gradient ascent search



(a) Top-3 Results from Dataset

(b) Top-3 Results from GAN





# Photos to paintings






# Paintings to photos







horse  $\rightarrow$  zebra



zebra  $\rightarrow$  horse



winter Yosemite  $\rightarrow$  summer Yosemite



summer Yosemite  $\rightarrow$  winter Yosemite



apple  $\rightarrow$  orange



orange  $\rightarrow$  apple

# High resolution image synthesis

## Large Scale GAN Training for High Fidelity Natural Image Synthesis



2018